

# Content Concierge – Backend Intelligence & Personalization Report

## 1. Overview

This project implements the **backend intelligence layer** for a Content Concierge system designed to deliver **contextual, educational, and compliant investment insights** inside an investing application.

The system does **not give advice**.

Instead, it explains *what the user is seeing, why it may be happening, and how investors often interpret similar situations*, tailored to:

- **Who the user is** (experience & behavior)
- **Where they are in the app** (placement)
- **Why the insight is being shown now** (trigger)
- **What they are looking at** (focus ticker, if any)

The result is a **placement-aware, archetype-driven, non-advisory insight engine** suitable for a regulated financial environment.

---

## 2. Core Concepts & Definitions

### 2.1 Archetype (Who the user is *right now*)

An **archetype** represents the user's current behavioral state, not a permanent label.

**Archetypes implemented:**

- **INACTIVE** – low engagement / returning after a gap
- **EVERYDAY** – active, non-advanced investor
- **ADVANCED** – experienced investor (flagged via experience level)
- *(Prospect is planned, not yet implemented)*

**How archetype is determined:**

```
If inactivity_flag == true → INACTIVE  
Else if investment_experience_level == "advanced" → ADVANCED  
Else → EVERYDAY
```

Archetypes control **what kind of insight is appropriate**, not the UI.

---

## 2.2 Tier (How much wealth context the user has)

A **tier** is derived from total investable assets and represents *scale*, not sophistication.

**Tier mapping:**

- UNDER\_250K
- FROM\_250K\_TO\_1M
- OVER\_1M

Used to:

- Shape tone
- Gate future sophistication
- Support business segmentation

Currently informational, but foundational for later prioritization.

---

## 2.3 Placement (Where the insight appears)

A **placement** represents the UI surface where an insight is embedded.

**Placements implemented:**

- **INVESTMENT\_DASHBOARD** – portfolio-level orientation
- **POSITIONS** – holding-level inspection
- **PERFORMANCE** – interpretation of returns and risk

Placement determines:

- Scope (portfolio vs ticker)
  - Bundle selection
  - Insight framing
- 

## 2.4 Trigger (Why the insight is shown now)

A **trigger** describes the user action or moment that caused the insight to be generated.

Examples:

- APP\_OPEN
- TAB\_VIEW
- HOVER\_TICKER
- DWELL\_NO\_ACTION (future)
- REPEAT\_VIEW (future)

Triggers do **not** change facts, but change *intent* and *timing*.

---

## 2.5 Focus Ticker (What the user is inspecting)

When the user is interacting with a specific holding (e.g., hovering over AAPL), the request includes:

```
"focus_ticker": "AAPL"
```

This enables:

- Ticker-scoped insights
- Provider filtering
- Contextual explanations

If absent, insights are portfolio-level.

---

### 3. The Insight Pipeline (End-to-End)

#### Step 1: Input Payload

The request includes:

- User profile
  - Wealth snapshot
  - Holdings
  - Goals
  - Activity summary
  - Preferences
  - **Request context** (placement, trigger, focus ticker)
- 

#### Step 2: Normalization

Raw data is converted into a **normalized context**, including:

- Derived archetype
- Derived tier
- Holdings count
- Top holdings
- Dividend profile
- Goal progress
- Inactivity signals

This ensures **all downstream logic works off a clean, stable abstraction**, not raw schemas.

---

### **Step 3: Provider Fetch (Grounding)**

External providers (e.g., market data, analyst context) are queried.

Outputs are:

- **Items** (neutral facts)
  - **Citations** (for attribution)
  - No opinions, no advice
- 

### **Step 4: Bundle Planning (Core Intelligence)**

This is the **heart of the system**.

A **bundle** is a structured container of facts meant for *one coherent insight*.

Each bundle has:

- **kind** – what type of insight it is
- **facts** – bullet-style truths the LLM may use
- **citations** – optional grounding sources

#### **What bundles do:**

- Decide *what* the LLM is allowed to say
  - Prevent hallucination
  - Enforce compliance boundaries
- 

## **4. Bundles Implemented (signals.py)**

### **4.1 Inactive Bundles**

**Purpose:** Re-orient returning users.

**Function:**

```
build_inactive_activation_signals()
```

**Facts include:**

- User is inactive
- Difference between funding vs investing
- Diversification concepts
- Goal framing

**Used on:** Dashboard only

---

## 4.2 Everyday Bundles

**Everyday – Performance**

```
build_everyday_performance_signals()
```

Explains:

- How investors interpret performance
- Benchmark framing (educational only)
- Goal context

**Everyday – Positions**

```
build_everyday_positions_signals()
```

Explains:

- What investors look for when reviewing holdings
  - News, earnings, macro context
  - Concentration awareness
-

## 4.3 Advanced Bundles

### Advanced – Performance

`build_advanced_performance_signals()`

Explains:

- Concentration signals
- Risk exposure
- Performance drivers

### Advanced – Positions

`build_advanced_positions_signals()`

Explains:

- Catalysts
- Volatility context
- Scenario thinking

---

## 4.4 Market Trend Bundle

`build_market_trend_signals()`

Aggregates:

- Analyst sentiment
- Price context
- Thematic signals

Used as a **secondary context layer**, not the primary explanation.

---

## 5. Facts (What the LLM is allowed to say)

**Facts are the most important safety construct in the system.**

They are:

- Short, declarative statements
- Derived from user data or providers
- Non-prescriptive
- Non-predictive

Example facts:

- “Largest holding represents ~86% of tracked holdings value.”
- “Recent prices ranged between X and Y.”
- “User is viewing ticker AAPL.”

The LLM:

- **May only speak using these facts**
  - Cannot invent numbers
  - Cannot give advice
- 

## 6. LLM Realize + Judge Pattern

Every insight passes through **two LLM steps**:

### 6.1 Realize

- Generates headline, explanation, personal relevance
- Uses facts only
- Educational tone enforced

### 6.2 Judge

- Independently classifies output as PASS / BLOCK
- Ensures no advice language
- Blocks imperatives or recommendations

This creates a **self-auditing system**.

---

## 7. Insight Types (kind\_to\_type)

Each bundle kind maps to an **InsightType** used by the UI:

Bundle Kind	InsightType
goal_portfolio	GOAL_PROGRESS
inactive_activation	PORTFOLIO_COMPOSITION
everyday_performance	PORTFOLIO_COMPOSITION
everyday_positions	PORTFOLIO_COMPOSITION
advanced_performanc e	PORTFOLIO_COMPOSITION
advanced_positions	PORTFOLIO_COMPOSITION
market_trend	MARKET_TREND
positions_ticker	MARKET_TREND

This separates **content intent** from **rendering logic**.

---

## 8. Scope & Priority

Each insight includes:

- **scope**
  - PORTFOLIO – whole account
  - TICKER – specific holding

- **priority**
    - Determines ordering when multiple insights exist
- 

## 9. What This System Achieves

- Fully **placement-aware** insight generation
  - **Behavior-driven personalization** (not static personas)
  - Strict **non-advisory compliance**
  - Deterministic, auditable outputs
  - Clear extension points for:
    - Ranking
    - Benchmark math
    - Advisor triggers
    - Experimentation
- 

## 10. What Is Intentionally Not Done (Yet)

- No buy/sell recommendations
- No ranking optimization
- No benchmark calculations
- No multimedia content
- No advisor escalation logic

These are **future layers**, not missing features.

---

## 11. Final Summary

This implementation delivers a **production-grade intelligence backbone** for a financial content concierge.

It converts raw financial data into **contextual, compliant, and placement-specific insights**, driven by user behavior and experience rather than generic rules.

It is designed to scale safely into more advanced capabilities without re-architecting the system.